

Accelerating machine learning through GPU and specialized processor unit.

PAPER

Bingsheng He, Wenbin Fang, Qiong Luo, Naga K. Govindaraju, and Tuyong Wang. Mars: A MapReduce Framework on Graphics Processors. in PACT, 2008.

M. C. Daz, F. A. Gonzlez, and R. Ramos-Pollan. Accelerating common machine learning algorithms through GPGPU symbolic computing. In IEEE Computing Colombian Conference, Bogota, Colombia, 2015

Chuntao Hong, Dehao Chen, Wenguang Chen, Weimin Zheng, and Haibo Lin. MapCG: Writing Parallel Program Portable between CPU and GPU. In PACT, 2010.

Rajat Raina, Anand Madhavan and Andrew Y. Ng. Large-scale deep unsupervised learning using graphics processors. In Proceeding ICML '09 Proceedings of the 26th Annual International Conference on Machine Learning, 2009.

Bryan Catanzaro, Narayanan Sundaram and Kurt Keutzer. Fast support vector machine training and classification on graphics processors. In Proceeding ICML '08 Proceedings of the 25th international conference on Machine learning, 2008.

Z. Chen, J. Wang, H. He and X. Huang. A fast deep learning system using GPU. In Circuits and Systems (ISCAS), IEEE International Symposium on, 2014.

J. Lu, S. Young, I. Arel and J. Holleman. A 1 TOPS/W Analog Deep Machine-Learning Engine With Floating-Gate Storage in 0.13 μ m CMOS. In IEEE Journal of Solid-State Circuits, 2015.

Peilong Li, Yan Luo, Ning Zhang and Yu Cao. HeteroSpark: A heterogeneous CPU/GPU Spark platform for machine learning algorithms. In Networking, Architecture and Storage (NAS), , IEEE International Conference, 2015.